

HW5-Coding

January 11, 2024

1 Homework 5: Convolutional neural network (30 points)

In this part, you need to implement and train a convolutional neural network on the CIFAR-10 dataset with PyTorch. ### What is PyTorch?

PyTorch is a system for executing dynamic computational graphs over Tensor objects that behave similarly as numpy ndarray. It comes with a powerful automatic differentiation engine that removes the need for manual back-propagation.

1.0.1 Why?

- Our code will now run on GPUs! Much faster training. When using a framework like PyTorch or TensorFlow you can harness the power of the GPU for your own custom neural network architectures without having to write CUDA code directly (which is beyond the scope of this class).
- We want you to be ready to use one of these frameworks for your project so you can experiment more efficiently than if you were writing every feature you want to use by hand.
- We want you to stand on the shoulders of giants! TensorFlow and PyTorch are both excellent frameworks that will make your lives a lot easier, and now that you understand their guts, you are free to use them :)
- We want you to be exposed to the sort of deep learning code you might run into in academia or industry. ## How can I learn PyTorch?

Justin Johnson has made an excellent [tutorial](#) for PyTorch.

You can also find the detailed [API doc](#) here. If you have other questions that are not addressed by the API docs, the [PyTorch forum](#) is a much better place to ask than StackOverflow.

Install PyTorch and Skorch.

```
[1]: !pip install -q torch skorch torchvision torchtext
```

```
[2]: import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
import skorch
import sklearn
```

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

1.1 0. Tensor Operations (5 points)

Tensor operations are important in deep learning models. In this part, you are required to get familiar to some common tensor operations in PyTorch.

1.1.1 1) Tensor squeezing, unsqueezing and viewing

Tensor squeezing, unsqueezing and viewing are important methods to change the dimension of a Tensor, and the corresponding functions are [torch.squeeze](#), [torch.unsqueeze](#) and [torch.Tensor.view](#). Please read the documents of the functions, and finish the following practice.

```
[3]: # x is a tensor with size being (3, 2)
x = torch.Tensor([[1, 2],
                  [3, 4],
                  [5, 6]])

print(x.shape)
# Add two new dimensions to x by using the function torch.unsqueeze, so that
# the size of x becomes (3, 1, 2, 1).
x = torch.unsqueeze(torch.unsqueeze(x, 1), 3)
print(x.shape)
# Remove the two dimensions just added by using the function torch.squeeze,
# and change the size of x back to (3, 2).
x = torch.squeeze(x)
print(x.shape)
# x is now a two-dimensional tensor, or in other words a matrix. Now use the
# function torch.Tensor.view and change x to a one-dimensional vector with
# size being (6).
x = x.view(6)
print(x.shape)
```

```
torch.Size([3, 2])
torch.Size([3, 1, 2, 1])
torch.Size([3, 2])
torch.Size([6])
```

1.1.2 2) Tensor concatenation and stack

Tensor concatenation and stack are operations to combine small tensors into big tensors. The corresponding functions are [torch.cat](#) and [torch.stack](#). Please read the documents of the functions, and finish the following practice.

```
[4]: # x is a tensor with size being (3, 2)
x = torch.Tensor([[1, 2], [3, 4], [5, 6]])

# y is a tensor with size being (3, 2)
```

```

y = torch.Tensor([[[-1, -2], [-3, -4], [-5, -6]]])

# Our goal is to generate a tensor z with size as (2, 3, 2), and z[0,:,:] = x,
↳ z[1,:,:] = y.

# Use torch.stack to generate such a z
z = torch.stack((x, y))
print(z[0,:,:])
# Use torch.cat and torch.unsqueeze to generate such a z
z = torch.cat((torch.unsqueeze(x, 0), torch.unsqueeze(y, 0)))
print(z[1,:,:])

```

```

tensor([[1., 2.],
        [3., 4.],
        [5., 6.]])
tensor([[[-1., -2.],
        [-3., -4.],
        [-5., -6.]]])

```

1.1.3 3) Tensor expansion

Tensor expansion is to expand a tensor into a larger tensor along singleton dimensions. The corresponding functions are `torch.Tensor.expand` and `torch.Tensor.expand_as`. Please read the documents of the functions, and finish the following practice.

```

[5]: # x is a tensor with size being (3)
x = torch.Tensor([1, 2, 3])

# Our goal is to generate a tensor z with size (2, 3), so that z[0,:,:] = x,
↳ z[1,:,:] = x.

# [TO DO]
# Change the size of x into (1, 3) by using torch.unsqueeze.
x = torch.unsqueeze(x, 0)
print(x.shape)

# [TO DO]
# Then expand the new tensor to the target tensor by using torch.Tensor.expand.
z = x.expand((2, 3))
print(z.shape)

```

```

torch.Size([1, 3])
torch.Size([2, 3])

```

1.1.4 4) Tensor reduction in a given dimension

In deep learning, we often need to compute the mean/sum/max/min value in a given dimension of a tensor. Please read the document of `torch.mean`, `torch.sum`, `torch.max`, `torch.min`, `torch.topk`, and finish the following practice.

```
[6]: # x is a random tensor with size being (10, 50)
x = torch.randn(10, 50)

# Compute the mean value for each row of x.
# You need to generate a tensor x_mean of size (10), and x_mean[k, :] is the
  ↳ mean value of the k-th row of x.
x_mean = torch.mean(x, dim=1)
print(x_mean[3, ])

# Compute the sum value for each row of x.
# You need to generate a tensor x_sum of size (10).
x_sum = torch.sum(x, dim=1)
print(x_sum.shape)

# Compute the max value for each row of x.
# You need to generate a tensor x_max of size (10).
x_max, x_max_indices = torch.max(x, dim=1)
print(x_max.shape)

# Compute the min value for each row of x.
# You need to generate a tensor x_min of size (10).
x_min, x_min_indices = torch.min(x, dim=1)
print(x_min.shape)

# Compute the top-5 values for each row of x.
# You need to generate a tensor x_top of size (10, 5), and x_top[k, :] is the
  ↳ top-5 values of each row in x.
x_xtop, x_xtop_indices = torch.topk(x, k=5, dim=1)
print(x_xtop.shape)

tensor(-0.0311)
torch.Size([10])
torch.Size([10])
torch.Size([10])
torch.Size([10, 5])
```

1.2 Convolutional Neural Networks

Implement a convolutional neural network for image classification on CIFAR-10 dataset.

CIFAR-10 is an image dataset of 10 categories. Each image has a size of 32x32 pixels. The following code will download the dataset, and split it into **train** and **test**. For this question, we use the default validation split generated by Skorch.

```
[7]: train = torchvision.datasets.CIFAR10("./data", train=True, download=True)
test = torchvision.datasets.CIFAR10("./data", train=False, download=True)
```

Files already downloaded and verified
Files already downloaded and verified

The following code visualizes some samples in the dataset. You may use it to debug your model if necessary.

```
[8]: def plot(data, labels=None, num_sample=5):
    n = min(len(data), num_sample)
    for i in range(n):
        plt.subplot(1, n, i+1)
        plt.imshow(data[i], cmap="gray")
        plt.xticks([])
        plt.yticks([])
        if labels is not None:
            plt.title(labels[i])

train.labels = [train.classes[target] for target in train.targets]
plot(train.data, train.labels)
```



1.2.1 1) Basic CNN implementation

Consider a basic CNN model

- It has 3 convolutional layers, followed by a linear layer.
- Each convolutional layer has a kernel size of 3, a padding of 1.
- ReLU activation is applied on every hidden layer.

Please implement this model in the following section. The hyperparameters is then be tuned and you need to fill the results in the table.

a) Implement convolutional layers (10 Points) Implement the initialization function and the forward function of the CNN.

```
[9]: class CNN(nn.Module):
    def __init__(self, channels):
        super(CNN, self).__init__()
        # implement parameter definitions here
        # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
        self.conv1 = nn.Conv2d(3, channels, kernel_size=3, padding=1)
        self.conv2 = nn.Conv2d(channels, channels, kernel_size=3, padding=1)
        self.conv3 = nn.Conv2d(channels, channels, kernel_size=3, padding=1)
        self.fc = nn.Linear(channels * 32 * 32, 10)
```

```

self.relu = nn.ReLU()
# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
def forward(self, images):
    # implement the forward function here
    # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
    images = self.relu(self.conv1(images))
    images = self.relu(self.conv2(images))
    images = self.relu(self.conv3(images))
    images = images.view(images.size(0), -1)
    images = self.fc(images)
    # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
    return images

```

b) **Tune hyperparameters** Train the CNN model on CIFAR-10 dataset. We can tune the number of channels, optimizer, learning rate and the number of epochs for best validation accuracy.

```

[10]: # implement hyperparameters, you can select and modify the hyperparameters by
      ↪yourself here.

optimize = [torch.optim.SGD, torch.optim.Adam]
learning_rate = [1e-3, 1e-2]
channel = [128, 256, 512]

train_data_normalized = torch.Tensor(train.data/255)
train_data_normalized = train_data_normalized.permute(0,3,1,2)

for l in learning_rate:
    for o in optimize:
        for c in channel:
            print(f'The channel was {c}, the learning rate was {l} and the optimizer
            ↪was {str(o)}')

            cnn = CNN(channels = c)

            model = skorch.NeuralNetClassifier(cnn, criterion=torch.nn.
            ↪CrossEntropyLoss,

                                                device="cuda",
                                                optimizer=o,
                                                lr=l,
                                                max_epochs=50,
                                                batch_size=32,
                                                callbacks=[skorch.callbacks.
            ↪EarlyStopping(lower_is_better=True)])

            # implement input normalization & type cast here
            model.fit(train_data_normalized, np.asarray(train.targets))

```

The channel was 128, the learning rate was 0.001 and the optimizer was <class 'torch.optim.sgd.SGD'>

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	2.1538	0.2984	1.9552	
2.6352				
2	1.8782	0.3644	1.8057	
2.5159				
3	1.7644	0.4022	1.7063	
2.5077				
4	1.6868	0.4274	1.6442	
2.5075				
5	1.6331	0.4393	1.5999	
2.5342				
6	1.5890	0.4508	1.5622	
2.5138				
7	1.5504	0.4631	1.5283	
2.5257				
8	1.5158	0.4754	1.4974	
2.4884				
9	1.4845	0.4845	1.4698	
2.5094				
10	1.4560	0.4921	1.4452	
2.5153				
11	1.4297	0.4997	1.4233	
2.5087				
12	1.4052	0.5050	1.4040	
2.5080				
13	1.3819	0.5106	1.3859	
2.4973				
14	1.3589	0.5151	1.3680	
2.5359				
15	1.3353	0.5195	1.3490	
2.5063				
16	1.3106	0.5262	1.3286	
2.5126				
17	1.2854	0.5358	1.3081	
2.5261				
18	1.2608	0.5433	1.2885	
2.5423				
19	1.2380	0.5490	1.2714	
2.5132				
20	1.2174	0.5529	1.2568	
2.5087				
21	1.1986	0.5548	1.2440	
2.5285				
22	1.1813	0.5578	1.2328	
2.5431				

23	1.1651	0.5624	1.2229
2.5441			
24	1.1498	0.5653	1.2139
2.5229			
25	1.1349	0.5673	1.2058
2.5197			
26	1.1205	0.5712	1.1980
2.5426			
27	1.1063	0.5731	1.1907
2.5121			
28	1.0922	0.5748	1.1839
2.5384			
29	1.0782	0.5791	1.1773
2.5581			
30	1.0641	0.5818	1.1709
2.5399			
31	1.0499	0.5848	1.1647
2.5395			
32	1.0356	0.5878	1.1586
2.5388			
33	1.0213	0.5893	1.1528
2.6760			
34	1.0068	0.5936	1.1473
2.5394			
35	0.9923	0.5969	1.1414
2.5441			
36	0.9779	0.5990	1.1361
2.5429			
37	0.9634	0.6004	1.1311
2.5244			
38	0.9489	0.6023	1.1261
2.5313			
39	0.9345	0.6055	1.1215
2.5420			
40	0.9201	0.6075	1.1168
2.5460			
41	0.9057	0.6098	1.1127
2.5793			
42	0.8912	0.6113	1.1087
2.5446			
43	0.8767	0.6123	1.1049
2.5421			
44	0.8621	0.6134	1.1013
2.5616			
45	0.8474	0.6151	1.0979
2.5381			
46	0.8326	0.6157	1.0951
2.5510			

47	0.8176	0.6175	1.0925
2.5599			
48	0.8026	0.6194	1.0902
2.5520			
49	0.7875	0.6209	1.0882
2.5809			
50	0.7722	0.6213	1.0867
2.5457			

The channel was 256, the learning rate was 0.001 and the optimizer was <class 'torch.optim.sgd.SGD'>

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	2.0563	0.3406	1.8789	
7.1793				
2	1.8134	0.3898	1.7401	
7.1448				
3	1.7079	0.4220	1.6553	
7.1275				
4	1.6390	0.4405	1.5975	
7.1238				
5	1.5826	0.4577	1.5463	
7.1552				
6	1.5311	0.4734	1.4980	
7.1274				
7	1.4837	0.4892	1.4555	
7.1010				
8	1.4430	0.5009	1.4227	
7.0573				
9	1.4095	0.5078	1.3980	
7.0830				
10	1.3811	0.5110	1.3779	
7.0372				
11	1.3549	0.5164	1.3582	
7.0741				
12	1.3289	0.5213	1.3379	
7.0458				
13	1.3020	0.5306	1.3161	
7.1066				
14	1.2740	0.5382	1.2939	
7.0709				
15	1.2460	0.5451	1.2722	
7.2062				
16	1.2196	0.5524	1.2528	
7.1100				
17	1.1956	0.5570	1.2357	
7.1131				
18	1.1735	0.5639	1.2209	
7.1158				

19	1.1529	0.5683	1.2078	
7.1141				
20	1.1333	0.5725	1.1960	
7.1172				
21	1.1144	0.5777	1.1853	
7.1421				
22	1.0959	0.5817	1.1752	
7.0767				
23	1.0775	0.5858	1.1655	
7.1139				
24	1.0591	0.5900	1.1564	
7.1460				
25	1.0407	0.5954	1.1477	
7.0990				
26	1.0222	0.5977	1.1392	
7.0885				
27	1.0036	0.6006	1.1308	
7.0724				
28	0.9851	0.6029	1.1229	
7.0989				
29	0.9666	0.6045	1.1153	
7.0775				
30	0.9482	0.6071	1.1081	
7.1163				
31	0.9300	0.6106	1.1012	
7.1367				
32	0.9118	0.6126	1.0948	
7.1540				
33	0.8938	0.6158	1.0887	
7.1020				
34	0.8758	0.6178	1.0832	
7.1123				
35	0.8579	0.6211	1.0781	
7.0862				
36	0.8400	0.6228	1.0733	
7.0589				
37	0.8221	0.6255	1.0692	
7.0586				
38	0.8041	0.6254	1.0656	7.0663
39	0.7862	0.6277	1.0625	
7.0586				
40	0.7682	0.6290	1.0601	
7.1094				
41	0.7503	0.6310	1.0586	
7.0598				
42	0.7325	0.6326	1.0580	
6.9988				
43	0.7148	0.6352	1.0584	7.0550

44	0.6971	0.6348	1.0597	7.0274
45	0.6796	0.6364	1.0621	7.0591
46	0.6622	0.6372	1.0657	7.0275

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 512, the learning rate was 0.001 and the optimizer was <class 'torch.optim.sgd.SGD'>

epoch	train_loss	valid_acc	valid_loss	dur
1	1.9754	0.3650	1.8107	
19.2832				
2	1.7411	0.4187	1.6621	
19.2475				
3	1.6337	0.4455	1.5773	
19.2669				
4	1.5530	0.4695	1.5058	
19.3221				
5	1.4831	0.4895	1.4450	
19.3939				
6	1.4225	0.5102	1.3954	
18.7571				
7	1.3679	0.5199	1.3532	
18.7726				
8	1.3187	0.5309	1.3183	
18.3529				
9	1.2777	0.5409	1.2904	
19.2355				
10	1.2442	0.5488	1.2685	
19.4449				
11	1.2156	0.5557	1.2501	
19.3896				
12	1.1897	0.5616	1.2341	
19.4934				
13	1.1651	0.5666	1.2193	
19.5106				
14	1.1410	0.5718	1.2054	
19.4175				
15	1.1172	0.5775	1.1923	
19.5110				
16	1.0936	0.5811	1.1802	
19.4205				
17	1.0702	0.5858	1.1688	
19.4421				
18	1.0472	0.5896	1.1582	
19.5591				
19	1.0246	0.5944	1.1484	
19.5666				
20	1.0022	0.5985	1.1392	
19.5558				

21	0.9801	0.6020	1.1307	
19.5756				
22	0.9582	0.6051	1.1225	
19.4179				
23	0.9364	0.6059	1.1148	
19.3417				
24	0.9146	0.6090	1.1076	
19.3634				
25	0.8930	0.6117	1.1009	
19.2752				
26	0.8716	0.6145	1.0952	
19.3319				
27	0.8503	0.6169	1.0905	
19.3732				
28	0.8294	0.6185	1.0869	
19.2923				
29	0.8087	0.6203	1.0845	
19.3363				
30	0.7882	0.6225	1.0834	
19.3682				
31	0.7679	0.6223	1.0835	19.3122
32	0.7478	0.6247	1.0848	19.3185
33	0.7277	0.6265	1.0873	19.3377
34	0.7077	0.6262	1.0910	19.3369

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 128, the learning rate was 0.001 and the optimizer was <class 'torch.optim.adam.Adam'>

epoch	train_loss	valid_acc	valid_loss	dur
1	1.5234	0.5851	1.1891	
2.6840				
2	1.0571	0.6169	1.0892	
2.6992				
3	0.8221	0.6141	1.1500	2.6971
4	0.6047	0.6024	1.3418	2.6915
5	0.4299	0.5875	1.6522	2.7243
6	0.3110	0.5717	2.0032	2.7339

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 256, the learning rate was 0.001 and the optimizer was <class 'torch.optim.adam.Adam'>

epoch	train_loss	valid_acc	valid_loss	dur
1	1.5872	0.5492	1.2596	
7.2608				
2	1.1153	0.6067	1.1385	
7.2793				
3	0.8639	0.6174	1.1836	7.2944
4	0.6562	0.5973	1.4399	7.3466

5	0.4779	0.5787	1.6401	7.3964
6	0.3543	0.5912	1.8194	7.3584

Stopping since valid_loss has not improved in the last 5 epochs.
The channel was 512, the learning rate was 0.001 and the optimizer was <class 'torch.optim.adam.Adam'>

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	2.0094	0.4524	1.5763	
20.1288				
2	1.3341	0.5933	1.1718	
20.0618				
3	1.0120	0.6028	1.1743	20.0439
4	0.7524	0.5812	1.4462	20.0632
5	0.5117	0.5493	1.8079	20.1398
6	0.3501	0.5658	2.0140	19.8778

Stopping since valid_loss has not improved in the last 5 epochs.
The channel was 128, the learning rate was 0.01 and the optimizer was <class 'torch.optim.sgd.SGD'>

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	1.7794	0.4516	1.5389	
2.5239				
2	1.4475	0.5071	1.3684	
2.5350				
3	1.3014	0.5455	1.2685	
2.5639				
4	1.1840	0.5665	1.2059	
2.5228				
5	1.0889	0.5884	1.1525	
2.5110				
6	0.9935	0.6106	1.0999	
2.5434				
7	0.9005	0.6267	1.0665	
2.4998				
8	0.8147	0.6328	1.0625	
2.5299				
9	0.7331	0.6365	1.0841	2.5514
10	0.6519	0.6339	1.1270	2.5452
11	0.5681	0.6330	1.1972	2.5207
12	0.4801	0.6275	1.3198	2.5524

Stopping since valid_loss has not improved in the last 5 epochs.
The channel was 256, the learning rate was 0.01 and the optimizer was <class 'torch.optim.sgd.SGD'>

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	1.7445	0.4624	1.4953	
7.0353				
2	1.3844	0.5303	1.3078	

7.0792				
3	1.2215	0.5637	1.2228	
7.0867				
4	1.0985	0.5894	1.1513	
7.0975				
5	0.9825	0.6145	1.0954	
7.0483				
6	0.8741	0.6258	1.0755	
7.0947				
7	0.7733	0.6318	1.0788	7.1635
8	0.6754	0.6339	1.1130	7.1180
9	0.5752	0.6342	1.1827	7.0919
10	0.4693	0.6280	1.2978	7.1209

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 512, the learning rate was 0.01 and the optimizer was <class 'torch.optim.sgd.SGD'>

epoch	train_loss	valid_acc	valid_loss	dur
1	1.7083	0.4843	1.4394	
19.4529				
2	1.3369	0.5464	1.2704	
19.4450				
3	1.1600	0.5789	1.1729	
19.5391				
4	1.0139	0.6071	1.1062	
19.3637				
5	0.8858	0.6250	1.0830	
19.4428				
6	0.7689	0.6330	1.0816	
19.3910				
7	0.6552	0.6359	1.1100	19.2039
8	0.5372	0.6317	1.1976	19.2425
9	0.4117	0.6256	1.3515	19.2913
10	0.2870	0.6163	1.5788	19.3785

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 128, the learning rate was 0.01 and the optimizer was <class 'torch.optim.adam.Adam'>

epoch	train_loss	valid_acc	valid_loss	dur
1	3.0769	0.1000	2.3039	
2.7074				
2	2.3042	0.1000	2.3039	2.6637
3	2.3042	0.1000	2.3039	2.6802
4	2.3042	0.1000	2.3039	2.6868
5	2.3042	0.1000	2.3039	2.7117

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 256, the learning rate was 0.01 and the optimizer was <class 'torch.optim.adam.Adam'>

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	6.5130	0.1000	2.3039	
7.2951				
2	2.3042	0.1000	2.3039	7.3146
3	2.3042	0.1000	2.3039	7.3648
4	2.3042	0.1000	2.3039	7.2656
5	2.3042	0.1000	2.3039	7.2715

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 512, the learning rate was 0.01 and the optimizer was <class 'torch.optim.adam.Adam'>

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	42.1812	0.1000	2.3039	
19.8690				
2	2.3042	0.1000	2.3039	19.7609
3	2.3042	0.1000	2.3039	19.7402
4	2.3042	0.1000	2.3039	19.7742
5	2.3042	0.1000	2.3039	19.8196

Stopping since valid_loss has not improved in the last 5 epochs.

Write down **validation accuracy** of your model under different hyperparameter settings. Note the validation set is automatically split by Skorch during `model.fit()`.

#channel for each layer	optimizer	SGD	Adam
128		0.6365	0.6169
256		0.6372	0.6147
512		0.6359	0.6028

1.2.2 2) Full CNN implementation (10 points)

Based on the CNN in the previous question, implement a full CNN model with max pooling layer.

- Add a max pooling layer after each convolutional layer.
- Each max pooling layer has a kernel size of 2 and a stride of 2.

Please implement this model in the following section. The hyperparameters is then be tuned and fill the results in the table. You are also required to complete the questions.

a) Implement max pooling layers Similar to the CNN implementation in previous question, implement max pooling layers.

```
[11]: class CNN_MaxPool(nn.Module):
    def __init__(self, channels):
        super(CNN_MaxPool, self).__init__()
        # implement parameter definitions here
        # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
        self.conv1 = nn.Conv2d(3, channels, kernel_size=3, padding=1)
        self.conv2 = nn.Conv2d(channels, channels, kernel_size=3, padding=1)
```

```

self.conv3 = nn.Conv2d(channels, channels, kernel_size=3, padding=1)
self.pool = nn.MaxPool2d(kernel_size=2, stride=2)
self.fc = nn.Linear(channels * 16, 10)
self.relu = nn.ReLU()
# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

def forward(self, images):
    # implement the forward function here
    # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
    images = self.pool(self.relu(self.conv1(images)))
    images = self.pool(self.relu(self.conv2(images)))
    images = self.pool(self.relu(self.conv3(images)))
    images = self.fc(images.view(images.size(0), -1))
    # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
    return images

```

b) **Tune hyperparameters** Based on the better optimizer found in the previous problem, we can tune the number of channels and learning rate for best validation accuracy.

```

[12]: # implement hyperparameters, you can select and modify the hyperparameters by
      ↪yourself here.
learning_rate = [1e-4, 1e-3, 1e-2]
channel = [128, 256, 512]
# Select the better optimizer by the result shown in the previous problem, you
      ↪can select and modify it by yourself here.
better_optimizer = torch.optim.SGD

train_data_normalized = torch.Tensor(train.data/255)
train_data_normalized = train_data_normalized.permute(0,3,1,2)

for l in learning_rate:
    for c in channel:
        print(f'The channel was {c}, the learning rate was {l}')

        cnn = CNN_MaxPool(channels = c)

        model = skorch.NeuralNetClassifier(cnn, criterion=torch.nn.
      ↪CrossEntropyLoss,

                                         device="cuda",
                                         optimizer=better_optimizer,
                                         optimizer__momentum=0.90,
                                         lr=l,
                                         max_epochs=500,
                                         batch_size=32,
                                         callbacks=[skorch.callbacks.
      ↪EarlyStopping(lower_is_better=True)])
        # implement input normalization & type cast here

```



```
model.fit(train_data_normalized, np.asarray(train.targets))
```

The channel was 128, the learning rate was 0.0001

epoch	train_loss	valid_acc	valid_loss	dur
1	2.2981	0.1008	2.2931	
1.2544				
2	2.2873	0.1160	2.2783	
1.2456				
3	2.2637	0.1913	2.2390	
1.2406				
4	2.1914	0.2647	2.1200	
1.2659				
5	2.0630	0.2773	2.0112	
1.2407				
6	1.9991	0.2887	1.9726	
1.2388				
7	1.9641	0.3045	1.9406	
1.2428				
8	1.9280	0.3187	1.9043	
1.2371				
9	1.8858	0.3386	1.8589	
1.2442				
10	1.8381	0.3637	1.8082	
1.2584				
11	1.7899	0.3852	1.7595	
1.2423				
12	1.7442	0.4009	1.7142	
1.2251				
13	1.7024	0.4138	1.6746	
1.2344				
14	1.6661	0.4223	1.6420	
1.2490				
15	1.6354	0.4299	1.6143	
1.2432				
16	1.6084	0.4401	1.5900	
1.2352				
17	1.5838	0.4467	1.5679	
1.2908				
18	1.5611	0.4516	1.5477	
1.2521				
19	1.5402	0.4585	1.5292	
1.2573				
20	1.5211	0.4652	1.5121	
1.2508				
21	1.5036	0.4728	1.4967	
1.2525				
22	1.4875	0.4788	1.4822	

1.2309			
23	1.4726	0.4852	1.4683
1.2362			
24	1.4586	0.4911	1.4552
1.2294			
25	1.4453	0.4969	1.4425
1.2294			
26	1.4325	0.5011	1.4299
1.2365			
27	1.4200	0.5061	1.4177
1.2458			
28	1.4078	0.5109	1.4060
1.2398			
29	1.3958	0.5138	1.3943
1.2350			
30	1.3839	0.5173	1.3820
1.2439			
31	1.3720	0.5226	1.3699
1.2308			
32	1.3602	0.5258	1.3580
1.2276			
33	1.3484	0.5283	1.3461
1.2491			
34	1.3367	0.5329	1.3342
1.2311			
35	1.3252	0.5367	1.3225
1.2333			
36	1.3139	0.5403	1.3114
1.2435			
37	1.3028	0.5437	1.3002
1.2431			
38	1.2918	0.5479	1.2895
1.2364			
39	1.2810	0.5519	1.2791
1.2408			
40	1.2705	0.5546	1.2688
1.2368			
41	1.2601	0.5595	1.2589
1.2404			
42	1.2499	0.5620	1.2493
1.2619			
43	1.2399	0.5665	1.2402
1.2810			
44	1.2301	0.5695	1.2310
1.2518			
45	1.2205	0.5723	1.2224
1.2632			
46	1.2111	0.5758	1.2136

1.2688			
47	1.2018	0.5788	1.2054
1.2762			
48	1.1927	0.5794	1.1972
1.3111			
49	1.1839	0.5825	1.1890
1.2514			
50	1.1752	0.5849	1.1816
1.2512			
51	1.1666	0.5883	1.1741
1.2772			
52	1.1583	0.5928	1.1668
1.2731			
53	1.1501	0.5940	1.1598
1.2501			
54	1.1420	0.5962	1.1530
1.2558			
55	1.1342	0.5984	1.1462
1.2485			
56	1.1264	0.6009	1.1398
1.2486			
57	1.1189	0.6030	1.1334
1.2641			
58	1.1116	0.6055	1.1269
1.2625			
59	1.1044	0.6076	1.1207
1.2528			
60	1.0973	0.6091	1.1148
1.2472			
61	1.0903	0.6111	1.1092
1.2445			
62	1.0835	0.6142	1.1034
1.2780			
63	1.0768	0.6174	1.0979
1.2460			
64	1.0702	0.6194	1.0930
1.2480			
65	1.0637	0.6210	1.0878
1.2386			
66	1.0573	0.6230	1.0830
1.2646			
67	1.0510	0.6253	1.0780
1.2486			
68	1.0448	0.6269	1.0734
1.2703			
69	1.0387	0.6283	1.0687
1.2576			
70	1.0326	0.6301	1.0642

1.2739			
71	1.0266	0.6322	1.0599
1.2543			
72	1.0206	0.6338	1.0559
1.2459			
73	1.0147	0.6339	1.0516
1.2536			
74	1.0089	0.6364	1.0477
1.2660			
75	1.0031	0.6377	1.0436
1.2520			
76	0.9974	0.6397	1.0397
1.2618			
77	0.9917	0.6412	1.0360
1.2456			
78	0.9861	0.6424	1.0323
1.2511			
79	0.9805	0.6440	1.0291
1.2561			
80	0.9750	0.6443	1.0255
1.2574			
81	0.9696	0.6457	1.0218
1.2423			
82	0.9641	0.6461	1.0180
1.2455			
83	0.9587	0.6481	1.0143
1.2495			
84	0.9534	0.6498	1.0108
1.2673			
85	0.9481	0.6505	1.0075
1.2664			
86	0.9428	0.6519	1.0042
1.2642			
87	0.9376	0.6537	1.0012
1.2540			
88	0.9325	0.6549	0.9985
1.2519			
89	0.9274	0.6565	0.9954
1.2496			
90	0.9224	0.6569	0.9925
1.2376			
91	0.9174	0.6578	0.9894
1.2434			
92	0.9124	0.6591	0.9864
1.2525			
93	0.9074	0.6598	0.9833
1.2669			
94	0.9025	0.6612	0.9807

1.2446				
95	0.8977	0.6627	0.9778	
1.2639				
96	0.8929	0.6631	0.9749	
1.2478				
97	0.8882	0.6639	0.9721	
1.2631				
98	0.8835	0.6635	0.9698	1.2525
99	0.8788	0.6652	0.9671	
1.2460				
100	0.8741	0.6660	0.9646	
1.2557				
101	0.8695	0.6667	0.9622	
1.2157				
102	0.8649	0.6670	0.9600	
1.2226				
103	0.8604	0.6680	0.9579	
1.2470				
104	0.8560	0.6685	0.9556	
1.2471				
105	0.8515	0.6696	0.9531	
1.2636				
106	0.8471	0.6708	0.9509	
1.2453				
107	0.8427	0.6719	0.9488	
1.2648				
108	0.8384	0.6730	0.9467	
1.2416				
109	0.8341	0.6739	0.9449	
1.2390				
110	0.8299	0.6743	0.9427	
1.2494				
111	0.8256	0.6754	0.9408	
1.2491				
112	0.8214	0.6762	0.9387	
1.3029				
113	0.8171	0.6774	0.9368	
1.2335				
114	0.8130	0.6783	0.9350	
1.2382				
115	0.8088	0.6788	0.9337	
1.2641				
116	0.8047	0.6790	0.9319	
1.2284				
117	0.8006	0.6799	0.9305	
1.2815				
118	0.7965	0.6797	0.9289	1.2423
119	0.7924	0.6804	0.9277	

1.2344				
120	0.7884	0.6814	0.9267	
1.2422				
121	0.7844	0.6825	0.9251	
1.2359				
122	0.7804	0.6836	0.9243	
1.2330				
123	0.7765	0.6839	0.9228	
1.2388				
124	0.7725	0.6838	0.9213	1.2402
125	0.7686	0.6843	0.9204	
1.2427				
126	0.7647	0.6852	0.9196	
1.2349				
127	0.7609	0.6853	0.9184	
1.2369				
128	0.7570	0.6855	0.9180	
1.2403				
129	0.7532	0.6864	0.9166	
1.2687				
130	0.7494	0.6865	0.9161	
1.2445				
131	0.7457	0.6863	0.9149	1.2484
132	0.7419	0.6866	0.9139	
1.2538				
133	0.7382	0.6870	0.9134	
1.2598				
134	0.7344	0.6867	0.9122	1.2445
135	0.7308	0.6870	0.9109	1.2340
136	0.7271	0.6877	0.9100	
1.2483				
137	0.7234	0.6882	0.9089	
1.2399				
138	0.7198	0.6890	0.9078	
1.2385				
139	0.7161	0.6893	0.9069	
1.2498				
140	0.7124	0.6896	0.9061	
1.2430				
141	0.7089	0.6899	0.9049	
1.2465				
142	0.7053	0.6908	0.9041	
1.2577				
143	0.7016	0.6909	0.9034	
1.2394				
144	0.6981	0.6913	0.9032	
1.2413				
145	0.6946	0.6922	0.9023	

1.2326				
146	0.6911	0.6923	0.9022	
1.2295				
147	0.6875	0.6926	0.9014	
1.2506				
148	0.6840	0.6930	0.9010	
1.2510				
149	0.6805	0.6928	0.9002	1.2589
150	0.6770	0.6929	0.9003	1.2646
151	0.6736	0.6928	0.8998	1.2367
152	0.6702	0.6926	0.9000	1.2473
153	0.6667	0.6926	0.9000	1.2560
154	0.6633	0.6929	0.8997	1.2398
155	0.6599	0.6928	0.8996	1.2459
156	0.6565	0.6930	0.8992	1.2429
157	0.6531	0.6935	0.8989	
1.2432				
158	0.6497	0.6941	0.8987	
1.2375				
159	0.6464	0.6939	0.8989	1.2657
160	0.6429	0.6943	0.8987	
1.2539				
161	0.6396	0.6945	0.8985	
1.2493				
162	0.6362	0.6949	0.8977	
1.2395				
163	0.6329	0.6952	0.8974	
1.2434				
164	0.6295	0.6964	0.8972	
1.2459				
165	0.6262	0.6963	0.8976	1.2451
166	0.6229	0.6966	0.8970	
1.2489				
167	0.6196	0.6969	0.8970	1.2390
168	0.6163	0.6969	0.8966	1.2390
169	0.6130	0.6972	0.8965	
1.2420				
170	0.6097	0.6971	0.8970	1.2501
171	0.6064	0.6972	0.8972	1.2455
172	0.6031	0.6973	0.8973	1.2587
173	0.5999	0.6977	0.8974	1.2335

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 256, the learning rate was 0.0001

epoch	train_loss	valid_acc	valid_loss	dur
1	2.2903	0.1550	2.2764	
2.0048				
2	2.2449	0.2557	2.1880	

2.0367			
3	2.0923	0.2833	2.0082
1.9728			
4	1.9878	0.2994	1.9576
1.9713			
5	1.9433	0.3160	1.9156
1.9831			
6	1.8947	0.3414	1.8631
1.9908			
7	1.8377	0.3682	1.8015
1.9663			
8	1.7765	0.3975	1.7378
1.9583			
9	1.7172	0.4138	1.6813
1.9663			
10	1.6679	0.4275	1.6381
1.9672			
11	1.6291	0.4359	1.6038
1.9414			
12	1.5966	0.4453	1.5742
1.9511			
13	1.5674	0.4545	1.5478
1.9501			
14	1.5405	0.4639	1.5238
1.9752			
15	1.5154	0.4712	1.5013
1.9437			
16	1.4917	0.4786	1.4799
1.9782			
17	1.4694	0.4865	1.4592
1.9980			
18	1.4482	0.4944	1.4391
1.9624			
19	1.4282	0.5017	1.4202
2.0186			
20	1.4094	0.5087	1.4026
2.1429			
21	1.3914	0.5141	1.3858
1.9936			
22	1.3742	0.5193	1.3698
1.9731			
23	1.3577	0.5244	1.3539
1.9545			
24	1.3417	0.5297	1.3381
1.9676			
25	1.3262	0.5351	1.3228
2.0107			
26	1.3112	0.5422	1.3077

2.0154			
27	1.2967	0.5489	1.2933
1.9544			
28	1.2826	0.5542	1.2793
1.9334			
29	1.2689	0.5590	1.2659
1.9792			
30	1.2555	0.5633	1.2532
1.9778			
31	1.2426	0.5684	1.2409
1.9689			
32	1.2300	0.5723	1.2291
1.9635			
33	1.2178	0.5752	1.2180
1.9777			
34	1.2059	0.5789	1.2072
1.9589			
35	1.1945	0.5826	1.1969
1.9786			
36	1.1833	0.5852	1.1869
1.9523			
37	1.1725	0.5895	1.1772
1.9688			
38	1.1621	0.5923	1.1680
1.9473			
39	1.1519	0.5945	1.1592
1.9135			
40	1.1420	0.5976	1.1508
1.9208			
41	1.1323	0.5991	1.1423
1.9732			
42	1.1229	0.6027	1.1343
1.9566			
43	1.1137	0.6052	1.1265
1.9981			
44	1.1047	0.6077	1.1189
2.0796			
45	1.0959	0.6095	1.1118
2.1442			
46	1.0872	0.6112	1.1049
2.0307			
47	1.0788	0.6140	1.0983
2.0678			
48	1.0705	0.6178	1.0917
2.0187			
49	1.0623	0.6202	1.0850
2.0403			
50	1.0543	0.6222	1.0789

2.0308			
51	1.0464	0.6250	1.0724
2.1660			
52	1.0387	0.6281	1.0665
2.0301			
53	1.0310	0.6291	1.0608
2.0476			
54	1.0235	0.6311	1.0552
2.0517			
55	1.0161	0.6332	1.0496
2.0480			
56	1.0088	0.6365	1.0441
2.0760			
57	1.0016	0.6385	1.0388
2.0316			
58	0.9944	0.6401	1.0335
2.0642			
59	0.9872	0.6420	1.0284
2.0728			
60	0.9802	0.6436	1.0236
2.0329			
61	0.9732	0.6458	1.0188
2.0265			
62	0.9663	0.6471	1.0141
2.0095			
63	0.9595	0.6491	1.0092
2.0621			
64	0.9527	0.6506	1.0044
2.0649			
65	0.9459	0.6520	0.9998
2.0479			
66	0.9392	0.6535	0.9956
2.0198			
67	0.9326	0.6545	0.9911
2.0393			
68	0.9260	0.6557	0.9865
2.0554			
69	0.9195	0.6580	0.9822
2.0564			
70	0.9130	0.6590	0.9779
2.0122			
71	0.9065	0.6601	0.9738
2.0433			
72	0.9001	0.6618	0.9698
2.0309			
73	0.8937	0.6626	0.9659
2.0264			
74	0.8874	0.6627	0.9620

2.0271			
75	0.8812	0.6650	0.9582
2.0270			
76	0.8749	0.6660	0.9548
2.0442			
77	0.8688	0.6680	0.9510
2.0258			
78	0.8627	0.6696	0.9471
2.0518			
79	0.8566	0.6712	0.9433
2.0563			
80	0.8507	0.6721	0.9396
2.0658			
81	0.8447	0.6743	0.9363
2.0697			
82	0.8389	0.6755	0.9327
2.0295			
83	0.8330	0.6771	0.9294
2.0627			
84	0.8273	0.6777	0.9257
2.0321			
85	0.8215	0.6795	0.9224
2.0290			
86	0.8158	0.6799	0.9191
2.0629			
87	0.8102	0.6807	0.9162
2.0277			
88	0.8045	0.6818	0.9127
2.0242			
89	0.7990	0.6822	0.9097
2.0278			
90	0.7935	0.6830	0.9068
2.0422			
91	0.7880	0.6843	0.9040
2.0201			
92	0.7827	0.6844	0.9011
2.0660			
93	0.7773	0.6862	0.8985
2.0853			
94	0.7720	0.6863	0.8961
2.0685			
95	0.7667	0.6868	0.8939
2.0514			
96	0.7614	0.6883	0.8914
2.0493			
97	0.7561	0.6895	0.8889
2.0512			
98	0.7509	0.6906	0.8866

2.0296			
99	0.7457	0.6918	0.8843
2.0263			
100	0.7406	0.6923	0.8823
2.0315			
101	0.7355	0.6933	0.8796
2.0452			
102	0.7304	0.6952	0.8776
2.0468			
103	0.7253	0.6960	0.8756
2.0622			
104	0.7203	0.6962	0.8738
2.0462			
105	0.7152	0.6964	0.8720
2.0219			
106	0.7103	0.6968	0.8703
2.0615			
107	0.7053	0.6980	0.8687
2.0403			
108	0.7004	0.6997	0.8666
2.0554			
109	0.6955	0.7003	0.8650
2.0276			
110	0.6906	0.7005	0.8635
2.0263			
111	0.6857	0.7008	0.8621
2.0260			
112	0.6808	0.7010	0.8612
2.0502			
113	0.6760	0.7015	0.8599
2.0579			
114	0.6712	0.7023	0.8586
2.0614			
115	0.6663	0.7028	0.8574
2.0420			
116	0.6615	0.7039	0.8558
2.0442			
117	0.6567	0.7045	0.8546
2.0396			
118	0.6520	0.7050	0.8534
2.0352			
119	0.6472	0.7057	0.8526
2.0395			
120	0.6424	0.7058	0.8516
2.0554			
121	0.6377	0.7061	0.8511
2.0780			
122	0.6329	0.7070	0.8501

2.0524				
123	0.6282	0.7073	0.8498	
2.0583				
124	0.6235	0.7080	0.8490	
2.0361				
125	0.6188	0.7081	0.8485	
2.0400				
126	0.6141	0.7086	0.8479	
2.0471				
127	0.6094	0.7095	0.8476	
2.0227				
128	0.6048	0.7098	0.8472	
2.1048				
129	0.6001	0.7093	0.8468	2.0418
130	0.5954	0.7100	0.8464	
2.0521				
131	0.5908	0.7102	0.8463	
2.0521				
132	0.5861	0.7097	0.8466	2.0439
133	0.5815	0.7097	0.8466	2.0447
134	0.5768	0.7098	0.8463	2.0570
135	0.5722	0.7106	0.8466	2.0503

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 512, the learning rate was 0.0001

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	2.2784	0.2172	2.2421	
5.8419				
2	2.1572	0.2795	2.0531	
5.8371				
3	2.0003	0.2974	1.9655	
5.8715				
4	1.9360	0.3221	1.9092	
5.8313				
5	1.8735	0.3526	1.8395	
5.8006				
6	1.8012	0.3877	1.7618	
5.8270				
7	1.7293	0.4130	1.6916	
5.8308				
8	1.6694	0.4266	1.6393	
5.8585				
9	1.6218	0.4404	1.5964	
5.8416				
10	1.5811	0.4551	1.5581	
5.8279				
11	1.5449	0.4642	1.5240	
5.8287				

12	1.5128	0.4762	1.4937
5.8505			
13	1.4844	0.4857	1.4668
5.8500			
14	1.4590	0.4940	1.4427
5.8015			
15	1.4358	0.5027	1.4208
5.8185			
16	1.4143	0.5089	1.4003
5.7775			
17	1.3938	0.5177	1.3807
5.7809			
18	1.3741	0.5274	1.3618
5.8301			
19	1.3551	0.5306	1.3440
5.8479			
20	1.3369	0.5375	1.3267
5.8261			
21	1.3193	0.5432	1.3101
5.8638			
22	1.3023	0.5485	1.2939
5.8397			
23	1.2859	0.5539	1.2783
5.8306			
24	1.2701	0.5607	1.2636
5.8306			
25	1.2547	0.5658	1.2493
5.8299			
26	1.2398	0.5690	1.2356
5.8520			
27	1.2255	0.5725	1.2227
5.8412			
28	1.2115	0.5774	1.2100
5.8202			
29	1.1980	0.5818	1.1977
5.8272			
30	1.1848	0.5849	1.1859
5.8512			
31	1.1721	0.5894	1.1747
5.8791			
32	1.1597	0.5937	1.1642
5.8637			
33	1.1477	0.5972	1.1539
5.8385			
34	1.1360	0.6017	1.1440
5.8708			
35	1.1246	0.6037	1.1344
5.8588			

36	1.1135	0.6068	1.1251
5.8406			
37	1.1027	0.6116	1.1162
5.8191			
38	1.0921	0.6140	1.1076
5.8183			
39	1.0818	0.6164	1.0994
5.8400			
40	1.0717	0.6200	1.0912
5.8710			
41	1.0618	0.6226	1.0833
5.8396			
42	1.0522	0.6244	1.0758
5.8304			
43	1.0427	0.6275	1.0682
5.8195			
44	1.0335	0.6293	1.0610
5.8364			
45	1.0244	0.6317	1.0538
5.8691			
46	1.0154	0.6333	1.0469
5.8540			
47	1.0066	0.6352	1.0401
5.8213			
48	0.9980	0.6393	1.0336
5.8194			
49	0.9894	0.6408	1.0274
5.8574			
50	0.9810	0.6424	1.0211
5.8138			
51	0.9727	0.6446	1.0148
5.8313			
52	0.9645	0.6471	1.0088
5.8242			
53	0.9564	0.6495	1.0032
5.8398			
54	0.9484	0.6513	0.9975
5.9281			
55	0.9405	0.6533	0.9920
5.8626			
56	0.9327	0.6554	0.9866
5.8507			
57	0.9249	0.6569	0.9812
5.8775			
58	0.9172	0.6593	0.9760
5.8589			
59	0.9096	0.6608	0.9710
5.8561			

60	0.9021	0.6627	0.9661
5.8470			
61	0.8945	0.6644	0.9611
5.8473			
62	0.8872	0.6662	0.9564
5.8372			
63	0.8798	0.6673	0.9519
5.8407			
64	0.8725	0.6689	0.9474
5.8130			
65	0.8653	0.6701	0.9432
5.7768			
66	0.8581	0.6720	0.9390
5.8231			
67	0.8509	0.6736	0.9348
5.8268			
68	0.8439	0.6755	0.9311
5.7574			
69	0.8368	0.6765	0.9273
5.7802			
70	0.8299	0.6775	0.9238
5.7815			
71	0.8229	0.6785	0.9204
5.7865			
72	0.8161	0.6788	0.9170
5.7975			
73	0.8092	0.6804	0.9135
5.8057			
74	0.8024	0.6806	0.9106
5.7843			
75	0.7957	0.6818	0.9073
5.7680			
76	0.7890	0.6827	0.9045
5.8120			
77	0.7823	0.6829	0.9018
5.8518			
78	0.7757	0.6840	0.8989
5.8271			
79	0.7691	0.6851	0.8959
5.8293			
80	0.7626	0.6856	0.8934
5.8416			
81	0.7560	0.6865	0.8910
5.8508			
82	0.7496	0.6891	0.8888
5.8732			
83	0.7431	0.6897	0.8866
5.8831			

84	0.7367	0.6901	0.8844	
5.8483				
85	0.7303	0.6913	0.8825	
5.8200				
86	0.7239	0.6924	0.8802	
5.8591				
87	0.7176	0.6934	0.8783	
5.8310				
88	0.7113	0.6947	0.8765	
5.8286				
89	0.7050	0.6945	0.8747	5.8146
90	0.6987	0.6947	0.8729	5.8159
91	0.6925	0.6953	0.8716	
5.8428				
92	0.6862	0.6960	0.8702	
5.8774				
93	0.6800	0.6963	0.8690	
5.8220				
94	0.6738	0.6973	0.8673	
5.8113				
95	0.6676	0.6977	0.8658	
5.8110				
96	0.6615	0.6979	0.8650	
5.8214				
97	0.6552	0.6987	0.8638	
5.8180				
98	0.6491	0.6998	0.8623	
5.8396				
99	0.6429	0.6998	0.8611	5.8334
100	0.6368	0.7001	0.8600	
5.8647				
101	0.6306	0.7005	0.8597	
5.8714				
102	0.6245	0.7008	0.8586	
5.8747				
103	0.6183	0.7022	0.8580	
5.8685				
104	0.6122	0.7036	0.8573	
5.8599				
105	0.6061	0.7038	0.8566	
5.8744				
106	0.6000	0.7045	0.8560	
5.8581				
107	0.5939	0.7043	0.8556	5.8937
108	0.5877	0.7047	0.8554	
5.8723				
109	0.5816	0.7055	0.8550	
5.8590				

110	0.5755	0.7062	0.8546	
5.8729				
111	0.5694	0.7072	0.8542	
5.8289				
112	0.5633	0.7078	0.8543	5.8658
113	0.5572	0.7084	0.8542	
5.8715				
114	0.5510	0.7092	0.8545	5.8243
115	0.5449	0.7101	0.8547	5.8232

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 128, the learning rate was 0.001

epoch	train_loss	valid_acc	valid_loss	dur
1	2.0880	0.3435	1.8522	
1.2581				
2	1.7094	0.4222	1.6121	
1.2551				
3	1.5293	0.4779	1.4590	
1.2461				
4	1.4131	0.5166	1.3561	
1.2641				
5	1.3185	0.5545	1.2578	
1.2645				
6	1.2342	0.5851	1.1778	
1.2813				
7	1.1595	0.6092	1.1174	
1.2463				
8	1.0932	0.6240	1.0684	
1.2586				
9	1.0339	0.6405	1.0294	
1.2445				
10	0.9815	0.6539	0.9903	
1.2962				
11	0.9356	0.6654	0.9619	
1.2490				
12	0.8951	0.6706	0.9392	
1.2587				
13	0.8587	0.6792	0.9227	
1.2544				
14	0.8253	0.6838	0.9088	
1.2355				
15	0.7949	0.6847	0.9001	
1.2700				
16	0.7662	0.6857	0.8936	
1.2282				
17	0.7392	0.6896	0.8875	
1.2647				
18	0.7137	0.6924	0.8883	1.2491

19	0.6892	0.6926	0.8917	1.2563
20	0.6657	0.6920	0.8973	1.2501
21	0.6426	0.6930	0.8992	1.2597

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 256, the learning rate was 0.001

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	2.0294	0.3651	1.7795	
2.0356				
2	1.6356	0.4571	1.5238	
2.0602				
3	1.4573	0.5088	1.3801	
2.0516				
4	1.3353	0.5480	1.2726	
2.0518				
5	1.2342	0.5860	1.1688	
2.0617				
6	1.1468	0.6158	1.0963	
2.0561				
7	1.0718	0.6362	1.0431	
2.0449				
8	1.0073	0.6504	0.9990	
2.0351				
9	0.9512	0.6636	0.9614	
2.0497				
10	0.9013	0.6760	0.9290	
2.0321				
11	0.8565	0.6846	0.9053	
2.0647				
12	0.8158	0.6907	0.8833	
2.0046				
13	0.7780	0.6980	0.8638	
2.0379				
14	0.7421	0.7061	0.8489	
2.0325				
15	0.7077	0.7086	0.8370	
2.0447				
16	0.6743	0.7126	0.8294	
2.0551				
17	0.6419	0.7142	0.8287	
2.0244				
18	0.6100	0.7141	0.8292	2.0451
19	0.5787	0.7144	0.8338	2.0301
20	0.5482	0.7154	0.8410	1.9690
21	0.5178	0.7146	0.8501	2.0364

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 512, the learning rate was 0.001

epoch	train_loss	valid_acc	valid_loss	dur
-------	------------	-----------	------------	-----

epoch	train_loss	valid_acc	valid_loss	dur
1	1.9412	0.4006	1.6886	
5.8159				
2	1.5561	0.4832	1.4561	
5.8241				
3	1.3814	0.5317	1.3202	
5.7983				
4	1.2504	0.5825	1.1793	
5.8271				
5	1.1435	0.6172	1.0891	
5.7973				
6	1.0565	0.6367	1.0347	
5.8330				
7	0.9844	0.6513	0.9924	
5.8321				
8	0.9226	0.6649	0.9565	
5.8528				
9	0.8678	0.6723	0.9307	
5.8526				
10	0.8175	0.6816	0.9064	
5.8205				
11	0.7708	0.6907	0.8843	
5.8476				
12	0.7261	0.6977	0.8707	
5.8333				
13	0.6835	0.7034	0.8592	
5.8509				
14	0.6414	0.7077	0.8510	
5.8561				
15	0.5999	0.7095	0.8512	5.8616
16	0.5585	0.7137	0.8485	
5.8777				
17	0.5169	0.7134	0.8533	5.8712
18	0.4755	0.7154	0.8610	5.9445
19	0.4344	0.7151	0.8835	5.8926
20	0.3936	0.7166	0.9002	5.9293

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 128, the learning rate was 0.01

epoch	train_loss	valid_acc	valid_loss	dur
1	1.6945	0.5278	1.3239	
1.2576				
2	1.2017	0.6201	1.0778	
1.2562				
3	0.9874	0.6588	0.9771	
1.2651				
4	0.8541	0.6843	0.9079	
1.2584				

5	0.7477	0.6944	0.9007	
1.2542				
6	0.6615	0.6895	0.9414	1.2701
7	0.5976	0.6989	0.9709	1.2705
8	0.5369	0.7013	0.9904	1.2862
9	0.4870	0.6831	1.1277	1.2548

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 256, the learning rate was 0.01

epoch	train_loss	valid_acc	valid_loss	dur
1	1.6353	0.5433	1.2842	
2.0251				
2	1.1358	0.6521	0.9950	
2.0334				
3	0.9150	0.6876	0.9009	
2.0270				
4	0.7748	0.7010	0.8741	
1.9882				
5	0.6588	0.7116	0.8673	
2.0708				
6	0.5528	0.7065	0.9097	2.0375
7	0.4568	0.6892	1.0338	1.9990
8	0.3879	0.6826	1.1266	2.0244
9	0.3369	0.6901	1.2241	2.0188

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 512, the learning rate was 0.01

epoch	train_loss	valid_acc	valid_loss	dur
1	1.6198	0.5535	1.2718	
5.9230				
2	1.0995	0.6567	0.9795	
5.8892				
3	0.8872	0.6842	0.8936	
5.8864				
4	0.7348	0.7061	0.8603	
5.8774				
5	0.6078	0.6981	0.9368	5.8177
6	0.4952	0.7072	0.9728	5.8441
7	0.3956	0.6971	1.1146	5.8306
8	0.3316	0.7080	1.1339	5.8365

Stopping since valid_loss has not improved in the last 5 epochs.

Write down the **validation accuracy** of the model under different hyperparameter settings.

#channel for each layer	validation accuracy
128	0.7013
256	0.7154
512	0.7166

#channel for each layer validation accuracy

For the best model you have, test it on the test set.

```
[14]: # implement the same input normalization & type cast here
train_data_normalized = torch.Tensor(train.data / 255).permute(0, 3, 1, 2)
test_data_normalized = torch.Tensor(test.data/255)
test_data_normalized = test_data_normalized.permute(0,3,1,2)

optimizer = torch.optim.SGD
cnn = CNN_MaxPool(512)

model = skorch.NeuralNetClassifier(cnn, criterion=torch.nn.CrossEntropyLoss,
                                  device="cuda",
                                  optimizer=optimizer,
                                  optimizer__momentum=0.90,
                                  lr=1e-3,
                                  max_epochs=1000,
                                  batch_size=32,
                                  callbacks=[skorch.callbacks.
↳EarlyStopping(lower_is_better=True)])
model.fit(train_data_normalized, np.asarray(train.targets))

test.predictions = model.predict(test_data_normalized)
sklearn.metrics.accuracy_score(test.targets, test.predictions)
```

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	1.9479	0.3974	1.6958	
5.8922				
2	1.5578	0.4850	1.4522	
5.8591				
3	1.3842	0.5254	1.3217	
5.8517				
4	1.2565	0.5781	1.1911	
5.8742				
5	1.1525	0.6112	1.1068	
5.8812				
6	1.0659	0.6318	1.0428	
5.9010				
7	0.9925	0.6515	0.9954	
5.8533				
8	0.9288	0.6682	0.9521	
5.8652				
9	0.8725	0.6780	0.9187	
5.8729				
10	0.8217	0.6885	0.8927	

5.8721				
11	0.7745	0.6961	0.8742	
5.8528				
12	0.7296	0.6975	0.8604	
5.8843				
13	0.6868	0.7001	0.8505	
5.8754				
14	0.6453	0.7029	0.8476	
5.8646				
15	0.6037	0.7040	0.8464	
5.8845				
16	0.5628	0.7057	0.8526	5.8861
17	0.5219	0.7060	0.8667	5.8865
18	0.4814	0.7047	0.8870	5.8785
19	0.4410	0.7042	0.9110	5.8638

Stopping since valid_loss has not improved in the last 5 epochs.

[14]: 0.704

How much **test accuracy** do you get? What can you conclude for the design of CNN structure and tuning of hyperparameters? (5 points)

Your Answer:

1. Adding pooling layers can improve computational speed and model accuracy.
2. Too high a learning rate may lead to not reaching the global optimum, but constantly oscillating.
3. The SGD optimizer is more perfect than the Adam optimizer.
4. If there are too few convolutional layers, it will lead to incomplete feature extraction; too many convolutional layers, on the other hand, can lead to worse results.